

ml_benz

August 1, 2023

```
[1]: # importing libraries
import pandas as pd
import numpy as np
import seaborn as sns

import matplotlib.pyplot as plt
%matplotlib inline
```

0.1 Data Exploration

```
[2]: # uploading the dataset
d1= pd.read_csv("train.csv")
```

```
[3]: d1.head(10)
```

```
[3]:   ID      y  X0 X1  X2 X3 X4 X5 X6 X8 ... X375 X376 X377 X378 X379 \
0   0  130.81   k  v   at  a  d  u  j  o ...    0    0    1    0    0
1   6   88.53   k  t   av  e  d  y  l  o ...    1    0    0    0    0
2   7   76.26  az  w    n  c  d  x  j  x ...    0    0    0    0    0
3   9   80.62  az  t    n  f  d  x  l  e ...    0    0    0    0    0
4  13   78.02  az  v    n  f  d  h  d  n ...    0    0    0    0    0
5  18   92.93   t  b    e  c  d  g  h  s ...    0    0    1    0    0
6  24  128.76  al  r    e  f  d  f  h  s ...    0    0    0    0    0
7  25   91.91   o  l   as  f  d  f  j  a ...    0    0    0    0    0
8  27  108.67   w  s   as  e  d  f  i  h ...    1    0    0    0    0
9  30  126.99   j  b   aq  c  d  f  a  e ...    0    0    1    0    0
```

```
      X380 X382 X383 X384 X385
0         0     0     0     0     0
1         0     0     0     0     0
2         0     1     0     0     0
3         0     0     0     0     0
4         0     0     0     0     0
5         0     0     0     0     0
6         0     0     0     0     0
7         0     0     0     0     0
```

```

8      0      0      0      0      0
9      0      0      0      0      0

```

[10 rows x 378 columns]

```
[4]: d1.tail(10)
```

```

[4]:      ID      y  X0  X1  X2  X3  X4  X5  X6  X8  ...  X375  X376  X377  X378  \
4199  8395  88.24   t  aa  ay  c  d  aa  l  o  ...    1    0    0    0
4200  8397  108.59  z  aa  e  c  d  aa  i  w  ...    1    0    0    0
4201  8399  107.39  w  v  t  d  d  aa  h  g  ...    0    1    0    0
4202  8402  123.34  ap  l  s  c  d  aa  d  r  ...    0    0    0    0
4203  8403   85.71  aq  s  as  c  d  aa  a  g  ...    1    0    0    0
4204  8405  107.39  ak  s  as  c  d  aa  d  q  ...    1    0    0    0
4205  8406  108.77   j  o  t  d  d  aa  h  h  ...    0    1    0    0
4206  8412  109.22  ak  v  r  a  d  aa  g  e  ...    0    0    1    0
4207  8415   87.48  al  r  e  f  d  aa  l  u  ...    0    0    0    0
4208  8417  110.85  z   r  ae  c  d  aa  g  w  ...    1    0    0    0

```

```

      X379  X380  X382  X383  X384  X385
4199     0     0     0     0     0     0
4200     0     0     0     0     0     0
4201     0     0     0     0     0     0
4202     0     0     0     0     0     0
4203     0     0     0     0     0     0
4204     0     0     0     0     0     0
4205     0     0     0     0     0     0
4206     0     0     0     0     0     0
4207     0     0     0     0     0     0
4208     0     0     0     0     0     0

```

[10 rows x 378 columns]

```
[5]: d1.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4209 entries, 0 to 4208
Columns: 378 entries, ID to X385
dtypes: float64(1), int64(369), object(8)
memory usage: 12.1+ MB

```

```

[6]: pd.options.display.float_format = '{:,.4f}'.format
var = d1.var()
v1 = var.reset_index()
v1.columns = ["id", "values"]
variance= v1.sort_values("values",ascending=1)
variance.head()

```

```
[6]:      id  values
      275  X289  0.0000
      315  X330  0.0000
      254  X268  0.0000
      332  X347  0.0000
      97   X107  0.0000
```

0.2 Data Processing

```
[7]: # We will remove the variables with variance 0 and
      # We will also remove id since it has a huge variance
      var = variance.loc[variance["values"] < 0, "id"]
      data1 = d1.drop(var, axis=1)
      data1.drop("ID", axis=1, inplace=True)
      data1.head()
```

```
[7]:      y  X0 X1 X2 X3 X4 X5 X6 X8 X10 ... X375 X376 X377 X378 X379 \
0 130.8100  k  v  at  a  d  u  j  o    0 ...    0    0    1    0    0
1  88.5300  k  t  av  e  d  y  l  o    0 ...    1    0    0    0    0
2  76.2600  az w   n  c  d  x  j  x    0 ...    0    0    0    0    0
3  80.6200  az t   n  f  d  x  l  e    0 ...    0    0    0    0    0
4  78.0200  az v   n  f  d  h  d  n    0 ...    0    0    0    0    0

      X380 X382 X383 X384 X385
0      0    0    0    0    0
1      0    0    0    0    0
2      0    1    0    0    0
3      0    0    0    0    0
4      0    0    0    0    0

[5 rows x 377 columns]
```

```
[8]: data1.head()
```

```
[8]:      y  X0 X1 X2 X3 X4 X5 X6 X8 X10 ... X375 X376 X377 X378 X379 \
0 130.8100  k  v  at  a  d  u  j  o    0 ...    0    0    1    0    0
1  88.5300  k  t  av  e  d  y  l  o    0 ...    1    0    0    0    0
2  76.2600  az w   n  c  d  x  j  x    0 ...    0    0    0    0    0
3  80.6200  az t   n  f  d  x  l  e    0 ...    0    0    0    0    0
4  78.0200  az v   n  f  d  h  d  n    0 ...    0    0    0    0    0

      X380 X382 X383 X384 X385
0      0    0    0    0    0
1      0    0    0    0    0
2      0    1    0    0    0
3      0    0    0    0    0
```

```
4      0      0      0      0      0
```

```
[5 rows x 377 columns]
```

```
[9]: data1.isnull().sum()
```

```
[9]: y      0
     X0      0
     X1      0
     X2      0
     X3      0
     ..
     X380    0
     X382    0
     X383    0
     X384    0
     X385    0
     Length: 377, dtype: int64
```

```
[10]: data1.isnull().any(axis=1)
```

```
[10]: 0      False
     1      False
     2      False
     3      False
     4      False
     ...
     4204   False
     4205   False
     4206   False
     4207   False
     4208   False
     Length: 4209, dtype: bool
```

```
[11]: data1.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4209 entries, 0 to 4208
Columns: 377 entries, y to X385
dtypes: float64(1), int64(368), object(8)
memory usage: 12.1+ MB
```

```
[ ]:
```

```
[12]: #we do not have any missing values
```

```
[13]: data1.describe()
```

```
[13]:
```

	y	X10	X11	X12	X13	X14	\
count	4,209.0000	4,209.0000	4,209.0000	4,209.0000	4,209.0000	4,209.0000	
mean	100.6693	0.0133	0.0000	0.0751	0.0580	0.4281	
std	12.6794	0.1146	0.0000	0.2635	0.2337	0.4949	
min	72.1100	0.0000	0.0000	0.0000	0.0000	0.0000	
25%	90.8200	0.0000	0.0000	0.0000	0.0000	0.0000	
50%	99.1500	0.0000	0.0000	0.0000	0.0000	0.0000	
75%	109.0100	0.0000	0.0000	0.0000	0.0000	1.0000	
max	265.3200	1.0000	0.0000	1.0000	1.0000	1.0000	

	X15	X16	X17	X18	...	X375	X376	\
count	4,209.0000	4,209.0000	4,209.0000	4,209.0000	...	4,209.0000	4,209.0000	
mean	0.0005	0.0026	0.0076	0.0078	...	0.3188	0.0573	
std	0.0218	0.0511	0.0869	0.0882	...	0.4661	0.2324	
min	0.0000	0.0000	0.0000	0.0000	...	0.0000	0.0000	
25%	0.0000	0.0000	0.0000	0.0000	...	0.0000	0.0000	
50%	0.0000	0.0000	0.0000	0.0000	...	0.0000	0.0000	
75%	0.0000	0.0000	0.0000	0.0000	...	1.0000	0.0000	
max	1.0000	1.0000	1.0000	1.0000	...	1.0000	1.0000	

	X377	X378	X379	X380	X382	X383	\
count	4,209.0000	4,209.0000	4,209.0000	4,209.0000	4,209.0000	4,209.0000	
mean	0.3148	0.0207	0.0095	0.0081	0.0076	0.0017	
std	0.4645	0.1423	0.0970	0.0895	0.0869	0.0408	
min	0.0000	0.0000	0.0000	0.0000	0.0000	0.0000	
25%	0.0000	0.0000	0.0000	0.0000	0.0000	0.0000	
50%	0.0000	0.0000	0.0000	0.0000	0.0000	0.0000	
75%	1.0000	0.0000	0.0000	0.0000	0.0000	0.0000	
max	1.0000	1.0000	1.0000	1.0000	1.0000	1.0000	

	X384	X385
count	4,209.0000	4,209.0000
mean	0.0005	0.0014
std	0.0218	0.0377
min	0.0000	0.0000
25%	0.0000	0.0000
50%	0.0000	0.0000
75%	0.0000	0.0000
max	1.0000	1.0000

[8 rows x 369 columns]

```
[14]: #going to apply the dataframe for better flexibility and intuitive way of
      ↪ storing.
      c = data1.corr().abs()
```

```
[15]: # unstack for index labeling
b = c.unstack()
```

```
[16]: s= pd.DataFrame(b)
s.reset_index(inplace = True)
s.head()
```

```
[16]:  level_0 level_1      0
0      y      y 1.0000
1      y    X10 0.0270
2      y    X11    nan
3      y    X12 0.0898
4      y    X13 0.0483
```

```
[17]: s["flag"] = np.where(s["level_0"] == s["level_1"], "same", "not same")
s.columns.values[2] = "corr"
s.head()
```

```
[17]:  level_0 level_1  corr      flag
0      y      y 1.0000    same
1      y    X10 0.0270  not same
2      y    X11    nan  not same
3      y    X12 0.0898  not same
4      y    X13 0.0483  not same
```

```
[18]: # Remove the variables with correlation more than .9
# .loc is the function for slicing the data and here we are using label based
↳ slicing.
# s.loc[s["flag"] != "same",]
name = s.loc[(s["corr"] > .9) & (s["flag"] != "same") , "level_1"]
```

```
[19]: # going to findout unique elements and elements are sorted in array format
final_name = name.unique()
final_name
```

```
[19]: array(['X251', 'X382', 'X215', 'X54', 'X76', 'X136', 'X162', 'X232',
        'X263', 'X272', 'X276', 'X279', 'X328', 'X35', 'X37', 'X39', 'X31',
        'X33', 'X302', 'X66', 'X111', 'X113', 'X134', 'X147', 'X198',
        'X222', 'X129', 'X61', 'X120', 'X102', 'X214', 'X239', 'X370',
        'X29', 'X137', 'X324', 'X248', 'X253', 'X385', 'X52', 'X172',
        'X216', 'X379', 'X48', 'X213', 'X84', 'X244', 'X101', 'X179',
        'X348', 'X71', 'X90', 'X94', 'X99', 'X122', 'X217', 'X242', 'X243',
        'X249', 'X320', 'X245', 'X88', 'X150', 'X363', 'X80', 'X98', 'X53',
        'X371', 'X199', 'X119', 'X311', 'X118', 'X227', 'X264', 'X130',
        'X49', 'X128', 'X58', 'X140', 'X146', 'X138', 'X158', 'X96',
        'X226', 'X326', 'X219', 'X360', 'X157', 'X156', 'X142', 'X62',
        'X250', 'X262', 'X266', 'X378', 'X187', 'X194', 'X362', 'X186',
```

```
'X238', 'X265', 'X112', 'X247', 'X205', 'X204', 'X368', 'X67',
'X19', 'X155', 'X152', 'X125', 'X229', 'X228', 'X254', 'X189',
'X364', 'X365', 'X89', 'X358', 'X202', 'X60', 'X178', 'X14',
'X230', 'X314', 'X184', 'X126', 'X296', 'X295', 'X299', 'X298',
'X44', 'X261', 'X346', 'X352', 'X367', 'X337', 'X334', 'X331',
'X246', 'X240', 'X208', 'X108', 'X185', 'X63', 'X17'], dtype=object)
```

```
[20]: # going to drop unique elements
data2 = data1.drop(final_name,axis=1)
data2.head()
```

```
[20]:
```

	y	X0	X1	X2	X3	X4	X5	X6	X8	X10	...	X369	X372	X373	X374	X375	\
0	130.8100	k	v	at	a	d	u	j	o	0	...	0	0	0	0	0	
1	88.5300	k	t	av	e	d	y	l	o	0	...	0	0	0	0	1	
2	76.2600	az	w	n	c	d	x	j	x	0	...	0	0	0	0	0	
3	80.6200	az	t	n	f	d	x	l	e	0	...	0	1	0	0	0	
4	78.0200	az	v	n	f	d	h	d	n	0	...	0	0	0	0	0	

	X376	X377	X380	X383	X384
0	0	1	0	0	0
1	0	0	0	0	0
2	0	0	0	0	0
3	0	0	0	0	0
4	0	0	0	0	0

[5 rows x 231 columns]

```
[21]: data2.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4209 entries, 0 to 4208
Columns: 231 entries, y to X384
dtypes: float64(1), int64(222), object(8)
memory usage: 7.4+ MB
```

```
[22]: data1.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4209 entries, 0 to 4208
Columns: 377 entries, y to X385
dtypes: float64(1), int64(368), object(8)
memory usage: 12.1+ MB
```

0.3 Apply Label Encoder

```
[23]: char = data2.select_dtypes(exclude='number')
char
```

```
[23]:      X0 X1  X2 X3 X4  X5 X6 X8
0      k v  at  a  d   u  j  o
1      k t  av  e  d   y  l  o
2     az w   n  c  d   x  j  x
3     az t   n  f  d   x  l  e
4     az v   n  f  d   h  d  n
...
4204   ak s  as  c  d  aa d  q
4205    j o  t  d  d  aa h  h
4206   ak v   r  a  d  aa g  e
4207   al r   e  f  d  aa l  u
4208    z r  ae  c  d  aa g  w
```

[4209 rows x 8 columns]

```
[24]: num = data2.select_dtypes(include='number')
num.describe()
```

```
[24]:
```

	y	X10	X11	X12	X13	X15	\
count	4,209.0000	4,209.0000	4,209.0000	4,209.0000	4,209.0000	4,209.0000	
mean	100.6693	0.0133	0.0000	0.0751	0.0580	0.0005	
std	12.6794	0.1146	0.0000	0.2635	0.2337	0.0218	
min	72.1100	0.0000	0.0000	0.0000	0.0000	0.0000	
25%	90.8200	0.0000	0.0000	0.0000	0.0000	0.0000	
50%	99.1500	0.0000	0.0000	0.0000	0.0000	0.0000	
75%	109.0100	0.0000	0.0000	0.0000	0.0000	0.0000	
max	265.3200	1.0000	0.0000	1.0000	1.0000	1.0000	

	X16	X18	X20	X21	...	X369	X372	\
count	4,209.0000	4,209.0000	4,209.0000	4,209.0000	...	4,209.0000	4,209.0000	
mean	0.0026	0.0078	0.1428	0.0026	...	0.0005	0.0005	
std	0.0511	0.0882	0.3499	0.0511	...	0.0218	0.0218	
min	0.0000	0.0000	0.0000	0.0000	...	0.0000	0.0000	
25%	0.0000	0.0000	0.0000	0.0000	...	0.0000	0.0000	
50%	0.0000	0.0000	0.0000	0.0000	...	0.0000	0.0000	
75%	0.0000	0.0000	0.0000	0.0000	...	0.0000	0.0000	
max	1.0000	1.0000	1.0000	1.0000	...	1.0000	1.0000	

	X373	X374	X375	X376	X377	X380	\
count	4,209.0000	4,209.0000	4,209.0000	4,209.0000	4,209.0000	4,209.0000	
mean	0.0192	0.2274	0.3188	0.0573	0.3148	0.0081	
std	0.1374	0.4192	0.4661	0.2324	0.4645	0.0895	

min	0.0000	0.0000	0.0000	0.0000	0.0000	0.0000
25%	0.0000	0.0000	0.0000	0.0000	0.0000	0.0000
50%	0.0000	0.0000	0.0000	0.0000	0.0000	0.0000
75%	0.0000	0.0000	1.0000	0.0000	1.0000	0.0000
max	1.0000	1.0000	1.0000	1.0000	1.0000	1.0000

	X383	X384
count	4,209.0000	4,209.0000
mean	0.0017	0.0005
std	0.0408	0.0218
min	0.0000	0.0000
25%	0.0000	0.0000
50%	0.0000	0.0000
75%	0.0000	0.0000
max	1.0000	1.0000

[8 rows x 223 columns]

```
[25]: char1 = pd.get_dummies(char.astype(str),drop_first=True)
char1.head()
```

```
[25]:
```

	X0_aa	X0_ab	X0_ac	X0_ad	X0_af	X0_ai	X0_aj	X0_ak	X0_al	X0_am	...	\
0	0	0	0	0	0	0	0	0	0	0	...	
1	0	0	0	0	0	0	0	0	0	0	...	
2	0	0	0	0	0	0	0	0	0	0	...	
3	0	0	0	0	0	0	0	0	0	0	...	
4	0	0	0	0	0	0	0	0	0	0	...	

	X8_p	X8_q	X8_r	X8_s	X8_t	X8_u	X8_v	X8_w	X8_x	X8_y
0	0	0	0	0	0	0	0	0	0	0
1	0	0	0	0	0	0	0	0	0	0
2	0	0	0	0	0	0	0	0	1	0
3	0	0	0	0	0	0	0	0	0	0
4	0	0	0	0	0	0	0	0	0	0

[5 rows x 187 columns]

```
[26]: # going to concatenating both objects(char1,num) in row wise
final_data = pd.concat([char1,num],axis=1)
final_data.head()
```

```
[26]:
```

	X0_aa	X0_ab	X0_ac	X0_ad	X0_af	X0_ai	X0_aj	X0_ak	X0_al	X0_am	...	\
0	0	0	0	0	0	0	0	0	0	0	...	
1	0	0	0	0	0	0	0	0	0	0	...	
2	0	0	0	0	0	0	0	0	0	0	...	
3	0	0	0	0	0	0	0	0	0	0	...	
4	0	0	0	0	0	0	0	0	0	0	...	

	X369	X372	X373	X374	X375	X376	X377	X380	X383	X384
0	0	0	0	0	0	0	1	0	0	0
1	0	0	0	0	1	0	0	0	0	0
2	0	0	0	0	0	0	0	0	0	0
3	0	1	0	0	0	0	0	0	0	0
4	0	0	0	0	0	0	0	0	0	0

[5 rows x 410 columns]

0.4 Model Development

```
[27]: # splitting data into independent and dependent features
X = final_data.drop("y",axis=1) #axis=1,means we are referring to columns(to_
    ↳drop)
y = final_data.loc[:, "y"]
```

```
[28]: # going to do 30,70 split
from sklearn.model_selection import train_test_split
X_test,X_train,y_test,y_train = train_test_split( X,y, test_size = 0.
    ↳3,random_state=42)
```

```
[29]: import xgboost as xg
```

```
[30]: xgr = xg.XGBRegressor(objective = 'reg:squarederror',n_estimators = 10, seed =_
    ↳42)
```

```
[31]: xgr.fit(X_train,y_train)
```

```
[31]: XGBRegressor(base_score=0.5, booster=None, colsample_bylevel=1,
    colsample_bynode=1, colsample_bytree=1, gamma=0, gpu_id=-1,
    importance_type='gain', interaction_constraints=None,
    learning_rate=0.300000012, max_delta_step=0, max_depth=6,
    min_child_weight=1, missing=nan, monotone_constraints=None,
    n_estimators=10, n_jobs=0, num_parallel_tree=1, random_state=42,
    reg_alpha=0, reg_lambda=1, scale_pos_weight=1, seed=42,
    subsample=1, tree_method=None, validate_parameters=False,
    verbosity=None)
```

```
[32]: ypredict = xgr.predict(X_test)
```

```
[33]: # Create file for the competition submission
```

```
[34]: d = pd.DataFrame()
d["y_test"] = y_test
```

```
d["ypredict"] = ypredict
d["mp"] = abs((d["y_test"]- d["ypredict"])/d["y_test"])
```

```
[35]: d.head()
```

```
[35]:      y_test  ypredict      mp
370   95.1300   90.4262  0.0494
3392  117.3600  108.3275  0.0770
2208  109.0100  108.3275  0.0063
3942   93.7700   87.1329  0.0708
1105  103.4100   92.0061  0.1103
```

```
[36]: #ROOT MEAN SQUARE
from sklearn.metrics import mean_squared_error
```

```
[ ]:
```

```
[37]: rmse = np.sqrt(mean_squared_error(y_test, ypredict))
      RSME=("RMSE: %f" % (rmse))
      print(RSME)
```

```
RMSE: 8.559700
```

```
[42]: #Accuracy
from sklearn.metrics import mean_squared_error, r2_score

# evaluate predictions
```

```
[43]: predictions = [round(value) for value in ypredict]
```

```
[44]: y_test = [95.1300, 117.3600, 109.0100, 93.7700, 103.4100]
      ypredict = [90.4262, 108.3275, 108.3275, 87.1329, 92.0061]

      # Calculate metrics
      rmse = np.sqrt(mean_squared_error(y_test, ypredict))
      r2 = r2_score(y_test, ypredict)

      # Print the metrics
      print("RMSE:", rmse)
      print("R-squared:", r2)
```

```
RMSE: 7.460263112786301
```

```
R-squared: 0.280786050169866
```

0.5 Model Evaluation

```
[45]: from sklearn.model_selection import KFold
      from sklearn.model_selection import cross_val_score
```

```
[46]: kfold = KFold(n_splits=50)
      results = cross_val_score(xgr, X_train, y_train, cv=kfold)
      AC=("Accuracy: %.2f%% (%.2f%%)" % (results.mean()*100, results.std()*100))
      print(AC)
```

Accuracy: 47.66% (26.77%)

```
[47]: from sklearn.model_selection import cross_val_score
      accuracies = cross_val_score(estimator=xgr,X = X_train, y = y_train, cv = 10)
      print("Accuracy: {:.2f} %".format(accuracies.mean()*100))
```

Accuracy: 43.42 %

```
[48]: xgr.get_params
```

```
[48]: <bound method XGBModel.get_params of XGBRegressor(base_score=0.5, booster=None,
      colsample_bylevel=1,
      colsample_bynode=1, colsample_bytree=1, gamma=0, gpu_id=-1,
      importance_type='gain', interaction_constraints=None,
      learning_rate=0.300000012, max_delta_step=0, max_depth=6,
      min_child_weight=1, missing=nan, monotone_constraints=None,
      n_estimators=10, n_jobs=0, num_parallel_tree=1, random_state=42,
      reg_alpha=0, reg_lambda=1, scale_pos_weight=1, seed=42,
      subsample=1, tree_method=None, validate_parameters=False,
      verbosity=None)>
```

```
[ ]: from sklearn.model_selection import GridSearchCV
```

```
[ ]: parameters= [{"learning_rate": (0.05, 0.10, 0.15),
      "max_depth": [ 3, 4, 5, 6, 8],
      "min_child_weight": [ 1, 3, 5, 7],
      "gamma": [ 0.0, 0.1, 0.2],
      "colsample_bytree": [ 0.3, 0.4],}]
```

```
[ ]: grid_search = GridSearchCV(estimator = xgr,param_grid =parameters,scoring =_
      ↪'accuracy',cv = 10, n_jobs = -1)
```

```
[ ]: grid_search.fit(X_train, y_train)
      best_accuracy = grid_search.best_score_
      best_accuracy
```

```
[ ]: best_parameters = grid_search.best_params_  
print("Best Accuracy: {:.2f} %".format(best_accuracy*100))  
print("Best Parameters:", best_parameters)
```

```
[ ]: print(RSME)
```

```
[ ]: print(AC)
```

1 the RSME score 8.5

2 the KFold accuracy 47.6%

end

```
[ ]:
```

```
[49]:
```

```
[ ]:
```

```
[52]:
```

```
[ ]:
```

```
[ ]:
```